

Softmax Classifier

Objective

- How to classify handwritten digits from the MNIST database by using Softmax classifier.

Table of Contents

In this lab, you will use a single layer Softmax to classify handwritten digits from the MNIST database.

- [Make some Data](#)
- [Build a Softmax Classifier](#)
- [Define Softmax, Criterion Function, Optimizer, and Train the Model](#)
- [Analyze Results](#)

Estimated Time Needed: **25 min**

Preparation

We'll need the following libraries

```
In [1]: # Import the libraries we need for this Lab

# Using the following line code to install the torchvision library
# !mamba install -y torchvision

!pip install torchvision==0.9.1 torch==1.8.1
import torch
import torch.nn as nn
import torchvision.transforms as transforms
import torchvision.datasets as dsets
import matplotlib.pyplot as plt
import numpy as np
```

```
ERROR: Ignored the following yanked versions: 0.1.6, 0.1.7, 0.1.8, 0.1.9, 0.2.0, 0.2.1, 0.2.2, 0.2.2.post2, 0.2.2.post3
ERROR: Could not find a version that satisfies the requirement torchvision==0.9.1 (from versions: 0.21.0, 0.22.0, 0.22.1)
ERROR: No matching distribution found for torchvision==0.9.1
Defaulting to user installation because normal site-packages is not writeable
```

Use the following function to plot out the parameters of the Softmax function:

In [2]: *# The function to plot parameters*

```
def PlotParameters(model):
    W = model.state_dict()['linear.weight'].data
    w_min = W.min().item()
    w_max = W.max().item()
    fig, axes = plt.subplots(2, 5)
    fig.subplots_adjust(hspace=0.01, wspace=0.1)
    for i, ax in enumerate(axes.flat):
        if i < 10:

            # Set the Label for the sub-plot.
            ax.set_xlabel("class: {0}".format(i))

            # Plot the image.
            ax.imshow(W[i, :].view(28, 28), vmin=w_min, vmax=w_max, cmap='seismic')

            ax.set_xticks([])
            ax.set_yticks([])

            # Ensure the plot is shown correctly with multiple plots
            # in a single Notebook cell.
    plt.show()
```

Use the following function to visualize the data:

In [3]: *# Plot the data*

```
def show_data(data_sample):
    plt.imshow(data_sample[0].numpy().reshape(28, 28), cmap='gray')
    plt.title('y = ' + str(data_sample[1]))
```

Make Some Data

Load the training dataset by setting the parameters `train` to `True` and convert it to a tensor by placing a transform object in the argument `transform`.

In [4]: *# Create and print the training dataset*

```
train_dataset = datasets.MNIST(root='./data', train=True, download=True, transform=tra
print("Print the training dataset:\n", train_dataset)
```

Print the training dataset:

```
Dataset MNIST
  Number of datapoints: 60000
  Root location: ./data
  Split: Train
  StandardTransform
  Transform: ToTensor()
```

Load the testing dataset and convert it to a tensor by placing a transform object in the argument `transform`.

In [5]: *# Create and print the validating dataset*

```
validation_dataset = datasets.MNIST(root='./data', download=True, transform=transforms)
print("Print the validating dataset:\n", validation_dataset)
```

Print the validating dataset:

```
Dataset MNIST
  Number of datapoints: 60000
  Root location: ./data
  Split: Train
  StandardTransform
```

Transform: ToTensor()

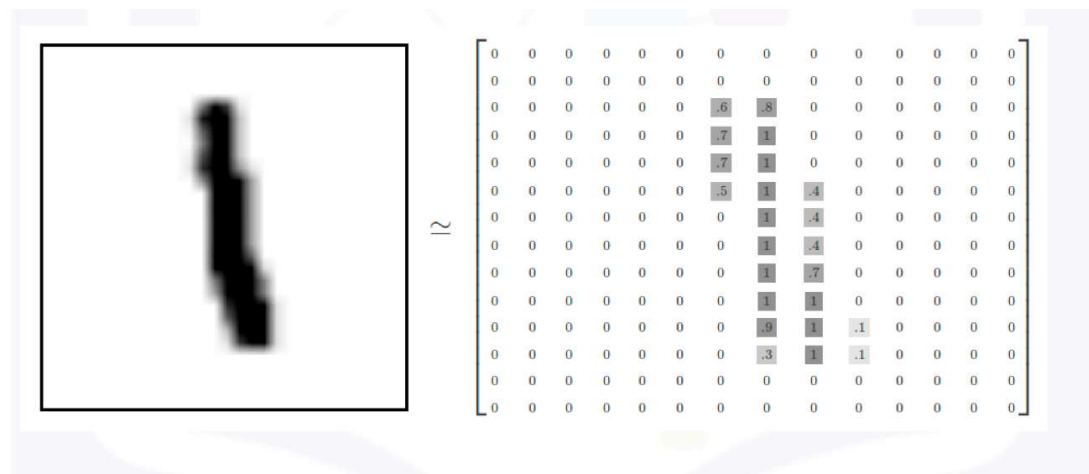
You can see that the data type is long:

In [6]: *# Print the type of the element*

```
print("Type of data element: ", type(train_dataset[0][1]))
```

Type of data element: <class 'int'>

Each element in the rectangular tensor corresponds to a number that represents a pixel intensity as demonstrated by the following image:



In this image, the values are inverted i.e back represents wight.

Print out the label of the fourth element:

In [7]: *# Print the Label*

```
print("The label: ", train_dataset[3][1])
```


The label: 1

The result shows the number in the image is 1

Plot the fourth sample:


```
In [8]: # Plot the image  
print("The image: ", show_data(train_dataset[3]))
```

The image: None



You see that it is a 1. Now, plot the third sample:

```
In [9]: # Plot the image  
show_data(train_dataset[2])
```



Build a Softmax Classifier

Build a Softmax classifier class:

```
In [10]: # Define softmax classifier class

class SoftMax(nn.Module):

    # Constructor
    def __init__(self, input_size, output_size):
        super(SoftMax, self).__init__()
        self.linear = nn.Linear(input_size, output_size)

    # Prediction
    def forward(self, x):
        z = self.linear(x)
        return z
```


The Softmax function requires vector inputs. Note that the vector shape is 28x28.

```
In [11]: # Print the shape of train dataset


train_dataset[0][0].shape
```

```
Out[11]: torch.Size([1, 28, 28])
```

Flatten the tensor as shown in this image:



The size of the tensor is now 784.



Set the input size and output size:

```
In [12]: # Set input size and output size

input_dim = 28 * 28
output_dim = 10
```

Define the Softmax Classifier, Criterion Function, Optimizer, and Train the Model

In [13]: *# Create the model*

```
model = SoftMax(input_dim, output_dim)
print("Print the model:\n ", model)
```

Print the model:

```
SoftMax(
  (linear): Linear(in_features=784, out_features=10, bias=True)
)
```

View the size of the model parameters:


In [14]: *# Print the parameters*

```
print('W: ', list(model.parameters())[0].size())
print('b: ', list(model.parameters())[1].size())
```

W: torch.Size([10, 784])

b: torch.Size([10])

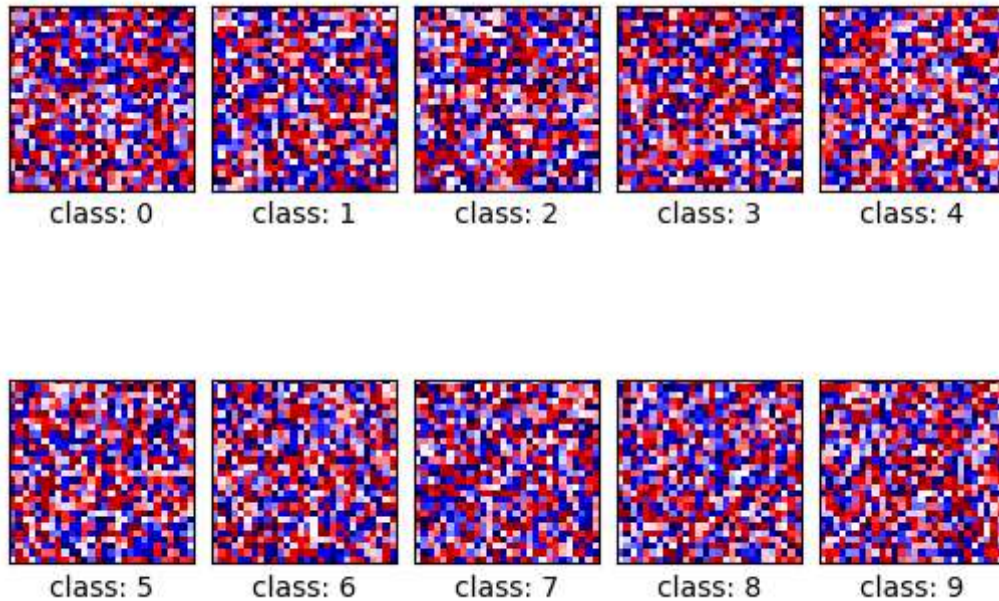
You can cover the model parameters for each class to a rectangular grid:



Plot the model parameters for each class as a square image:

In [15]: *# Plot the model parameters for each class*

```
PlotParameters(model)
```



Define the learning rate, optimizer, criterion, data loader:

```
In [16]: # Define the learning rate, optimizer, criterion and data loader

learning_rate = 0.1
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
criterion = nn.CrossEntropyLoss()
train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=100)
validation_loader = torch.utils.data.DataLoader(dataset=validation_dataset, batch_s
```

Train the model and determine validation accuracy **(should take a few minutes)**:

```
In [17]: # Train the model

n_epochs = 10
loss_list = []
accuracy_list = []
N_test = len(validation_dataset)

def train_model(n_epochs):
    for epoch in range(n_epochs):
        for x, y in train_loader:
            optimizer.zero_grad()
            z = model(x.view(-1, 28 * 28))
            loss = criterion(z, y)
            loss.backward()
            optimizer.step()

        correct = 0
        # perform a prediction on the validation data
        for x_test, y_test in validation_loader:
            z = model(x_test.view(-1, 28 * 28))
            _, yhat = torch.max(z.data, 1)
            correct += (yhat == y_test).sum().item()
        accuracy = correct / N_test
```



```

loss_list.append(loss.data)
accuracy_list.append(accuracy)

train_model(n_epochs)

```

Analyze Results

Plot the loss and accuracy on the validation data:


```

In [18]: # Plot the Loss and accuracy

fig, ax1 = plt.subplots()
color = 'tab:red'
ax1.plot(loss_list,color=color)
ax1.set_xlabel('epoch',color=color)
ax1.set_ylabel('total loss',color=color)
ax1.tick_params(axis='y', color=color)

ax2 = ax1.twinx()
color = 'tab:blue'
ax2.set_ylabel('accuracy', color=color)
ax2.plot( accuracy_list, color=color)
ax2.tick_params(axis='y', color=color)
fig.tight_layout()


```



View the results of the parameters for each class after the training. You can see that they look like the corresponding numbers.

In [19]: *# Plot the parameters*


```
PlotParameters(model)
```




We Plot the first five misclassified samples and the probability of that class.

In [20]: *# Plot the misclassified samples*


```
Softmax_fn=nn.Softmax(dim=-1)
count = 0
for x, y in validation_dataset:
    z = model(x.reshape(-1, 28 * 28))
    _, yhat = torch.max(z, 1)
    if yhat != y:
        show_data((x, y))
        plt.show()
        print("yhat:", yhat)
        print("probability of class ", torch.max(Softmax_fn(z)).item())
        count += 1
    if count >= 5:
        break
```




yhat: tensor([7])
probability of class 0.6424345374107361




yhat: tensor([7])
probability of class 0.7160860300064087



yhat: tensor([5])
probability of class 0.7200515270233154




yhat: tensor([8])
probability of class 0.37033313512802124




yhat: tensor([6])
 probability of class 0.5202871561050415

We Plot the first five correctly classified samples and the probability of that class, we see the probability is much larger.


```
In [21]: # Plot the classified samples
Softmax_fn=nn.Softmax(dim=-1)
count = 0
for x, y in validation_dataset:
    z = model(x.reshape(-1, 28 * 28))
    _, yhat = torch.max(z, 1)
    if yhat == y:
        show_data((x, y))
        plt.show()
        print("yhat:", yhat)
        print("probability of class ", torch.max(Softmax_fn(z)).item())
        count += 1
    if count >= 5:
        break
```




yhat: tensor([5])
probability of class 0.8634554743766785




yhat: tensor([0])
probability of class 0.9996826648712158



yhat: tensor([4])
probability of class 0.8717507123947144



yhat: tensor([1])
probability of class 0.9656764268875122



yhat: tensor([9])
probability of class 0.9250146150588989